

Forecasting in R

The forecasters' toolbox

Bahman Rostami-Tabar



Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

Learning outcome


You should be able to:

- 1 Discuss general tools that are useful for many different forecasting situations
- 2 Explain simple forecasting methods (benchmarks)
- 3 Specify and estimate models using R functions in fable
- 4 Recognise and extract fitted values and residuals
- 5 Produce point and prediction interval forecasts

Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

A tidy forecasting workflow




Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation**
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

Data preparation and visualisation

```
# Set training data from 1992 to 2007  
train <- aus_production %>%  
  filter(between(year(Quarter), 1992, 2007))  
train <- aus_production %>%  
  filter_index(1992 ~ 2007)  
train %>% autoplot(Beer)
```




Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)**
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

Some simple forecasting methods

MEAN(y): Average method


- Forecast of all future values is equal to mean of historical data $\{y_1, \dots, y_T\}$.
- Forecasts: $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$



Some simple forecasting methods

NAIVE(y): Naïve method


- Forecasts equal to last observed value.
- Forecasts: $\hat{y}_{T+h|T} = y_T$.
- Consequence of efficient market hypothesis.



Some simple forecasting methods

SNAIVE ($y \sim \text{lag}(m)$): Seasonal naïve method

- Forecasts equal to last value from same season.
- Forecasts: $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$, where m = seasonal period and k is the integer part of $(h - 1)/m$.



Some simple forecasting methods

RW(y ~ drift()): Drift method


- Forecasts equal to last value plus average change.
- Forecasts:

$$\begin{aligned}\hat{y}_{T+h|T} &= y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) \\ &= y_T + h \left(\frac{y_T - y_1}{T-1} \right).\end{aligned}$$

- Equivalent to extrapolating a line drawn between first and last observations.

Some simple forecasting methods

Drift method



Model specification

- Model specification in fable supports a formula based interface
- A model formula in R is expressed using `response ~ terms`
 - ▶ the formula's left side describes the response
 - ▶ the right describes terms used to model the response.
- Attention: `MODEL_NAME` is in capital letters, e.g. `SNAIVE`

```
MODEL_NAME(response_variable ~ term1+term2+...)
```

```
SNAIVE(Beer ~ lag("year"))
```

Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)**
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals

Model estimation: template

The `model()` function trains models to data. - It returns a model table or a mable object.

```
# Fit the models  
my_mable <- my_data %>%  
  model(  
    choose_name1 = MODEL_1(response_variable ~ term1+...),  
    choose_name2 = MODEL_2(response_variable ~ term1+...),  
    choose_name3 = MODEL_3(response_variable ~ term1+...),  
    choose_name4 = MODEL_4(response_variable ~ term1+...)  
  )
```

Model estimation

```
# Fit the models
```

```
beer_fit <- train %>%
```

```
  model(
```

```
    mean = MEAN(Beer),
```

```
    naive = NAIVE(Beer),
```

```
    snaive = SNAIVE(Beer, lag="year"),
```

```
    drift = RW(Beer ~ drift())
```

```
)
```

```
#beer_fit <- beer_fit %>% stream(new_data),
```

```
#we can update the fitted models once we have new data
```

mable: a model object

```
beer_fit
```

```
## # A mable: 1 x 4
##   mean    naive    snaive    drift
##   <model> <model> <model>   <model>
## 1 <MEAN>  <NAIVE> <SNAIVE> <RW w/ drift>
```

- A mable is a model table, each cell corresponds to a fitted model.
- A mable contains
 - ▶ a row for each time series
 - ▶ a column for each model specification

Extract coefficients from mab1e

```
beer_fit %>% select(snaive) %>% report()  
beer_fit %>% tidy()  
beer_fit %>% glance()
```

- The `report()` function gives a formatted model-specific display.
- The `tidy()` function is used to extract the coefficients from the models.
- The `glance()` shows a summary from the models.
- We can extract information about some specific model using the `filter()` and `select()` functions.

Producing forecasts

- The `forecast()` function is used to produce forecasts from estimated models.
- **h** can be specified with a number (the number of future observations) or natural language (the length of time to predict).

```
beer_fc <- beer_fit %>%  
  forecast(h = "3 years")  
#h = "3 years" is equivalent to setting h = 12.
```

Producing forecasts

```
## # A fable: 48 x 4 [1Q]
## # Key:      .model [4]
##   .model Quarter Beer .distribution
##   <chr>      <qtr> <dbl> <dist>
## 1 mean      2007 Q2  436. N(436, 1963)
## 2 mean      2007 Q3  436. N(436, 1963)
## 3 mean      2007 Q4  436. N(436, 1963)
## 4 mean      2008 Q1  436. N(436, 1963)
## # ... with 44 more rows
```


A fable is a forecast table with point forecasts and distributions.

Check model performance

Once a model has been fitted, it is important to check how well it has performed on the data. I come back to this latter.

Visualising forecasts

```
# Plot forecasts against actual values
beer_fc %>%
  autoplot(train, level = NULL) +
  autolayer(filter_index(aus_production, "2007 Q1" ~ .), color = "black") +
  ggtitle("Forecasts for quarterly beer production") +
  xlab("Year") + ylab("Megalitres") +
  guides(colour=guide_legend(title="Forecast"))
```



Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preperation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals**
- 7 lab session 5
- 8 Prediction intervals

Fitted values


- $\hat{y}_{T|T-1}$ is the forecast of y_T based on observations y_1, \dots, y_{T-1} .
- We call these “fitted values”.
- Sometimes drop the subscript: $\hat{y}_T \equiv \hat{y}_{T|T-1}$.
- Often not true forecasts since parameters are estimated on all data.

For example:

- $\hat{y}_T = \bar{y}$ for average method.
- $\hat{y}_T = y_{T-1} + (y_T - y_1)/(T - 1)$ for drift method.
- $\hat{y}_T = y_{T-1}$ for naive method.

Fitted values

```
beer_fit %>% select(naive) %>% augment() %>%  
  ggplot(aes(x=Quarter, y=Beer))+  
  geom_line()+  
  geom_line(aes(y=.fitted), colour="red")
```




Residuals

- The “residuals” in a time series model are what is left over after fitting a model.
- Residuals are useful in checking whether a model has adequately captured the information in the data.

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Residuals

```
#beer_fit %>% fitted  
#augment() function gets residuals and fitted values  
beer_fit %>% select(snaive) %>% augment() %>%  
  ggplot(aes(x=Quarter, y=.resid))+  
  geom_line()
```



Extract fitted values and residuals

```
beer_fit %>% augment()
```

```
## # A tsibble: 244 x 5 [1Q]
## # Key:           .model [4]
##   .model Quarter Beer .fitted .resid
##   <chr>      <qtr> <dbl>   <dbl> <dbl>
## 1 mean      1992 Q1   443     436.   6.70
## 2 mean      1992 Q2   410     436. -26.3
## 3 mean      1992 Q3   420     436. -16.3
## 4 mean      1992 Q4   532     436.  95.7
## 5 mean      1993 Q1   433     436.  -3.30
## 6 mean      1993 Q2   421     436. -15.3
## 7 mean      1993 Q3   410     436. -26.3
## 8 mean      1993 Q4   512     436.  75.7
## 9 mean      1994 Q1   449     436.  12.7
## 10 mean     1994 Q2   381     436. -55.3
```

Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5**
- 8 Prediction intervals

lab session 5


- Produce forecasts from the bechmark methods for daily A&E series for 42 days
- Plot the results using `autoplot()`.
 - ▶ Use `filter_index()` to show the plot from 2016
- Use `augment()` to extract fitted values for snaive method
- Extract residuals for mean method

Outline

- 1 Learning outcome
- 2 A tidy forecasting workflow
- 3 Data preparation(tidy) and visualisation
- 4 Define the model (specify)
- 5 Train the model (estimate)
- 6 Fitted values and Residuals
- 7 lab session 5
- 8 Prediction intervals**

Importance of providing interval forecast


Point forecasts are often useless without a measure of uncertainty



"Yes sir, you can absolutely trust those numbers"

Prediction intervals

- A prediction interval gives a region within which we expect y_{T+h} to lie with a specified probability
- It consists of an upper and a lower limit between which the future value is expected to lie



Prediction intervals

- Assuming forecast errors are normally distributed, then a c% PI is:

$$\hat{y}_{T+h|T} \pm c\hat{\sigma}_h$$

where the multiplier c depends on the coverage probability and $\hat{\sigma}_h$ is the st dev of the h -step distribution.

Prediction intervals

- Forecast intervals can be extracted using the `hilo()` function
- Use `level` argument to control coverage.

```
fit <- train %>% model(NAIVE(Beer))  
forecast(fit) %>% hilo(level = c(80, 95))
```

```
## # A tsibble: 8 x 5 [1Q]  
## # Key:           .model [1]  
##   .model      Quarter Beer      80%      95%  
##   <chr>        <qtr> <dbl>    <hilo>    <hilo>  
## 1 NAIVE(Be~  2007 Q2   427 [342.5627, 511~ [297.86430, 556~  
## 2 NAIVE(Be~  2007 Q3   427 [307.5876, 546~ [244.37454, 609~  
## 3 NAIVE(Be~  2007 Q4   427 [280.7503, 573~ [203.33041, 650~37
```